

# LA-UR-23-32613

Approved for public release; distribution is unlimited.

**Title:** Implicit Monte Carlo at LANL-Progress and Challenges

**Author(s):** Long, Alex Roberts

**Intended for:** Monte Carlo Computational Summit, 2023-10-25/2023-10-26 (South Bend, Indiana, United States)

**Issued:** 2023-11-06



Los Alamos National Laboratory, an affirmative action/equal opportunity employer, is operated by Triad National Security, LLC for the National Nuclear Security Administration of U.S. Department of Energy under contract 89233218CNA000001. By approving this article, the publisher recognizes that the U.S. Government retains nonexclusive, royalty-free license to publish or reproduce the published form of this contribution, or to allow others to do so, for U.S. Government purposes. Los Alamos National Laboratory requests that the publisher identify this article as work performed under the auspices of the U.S. Department of Energy. Los Alamos National Laboratory strongly supports academic freedom and a researcher's right to publish; as an institution, however, the Laboratory does not endorse the viewpoint of a publication or guarantee its technical correctness.

# Implicit Monte Carlo at LANL— Progress and Challenges

Alex Long, Los Alamos National Laboratory

Jayenne IMC Team: Matt Cleveland, Ryan Wollaeger,  
Kendra Long, Ben Ryan, Irina Sagert, Daniel Holladay

10/25/2023

LA-UR-2023-XXXXX

# The TRT equations are non-linear in temperature

- The Thermal Radiative Transfer (TRT) equations, without physical scattering:

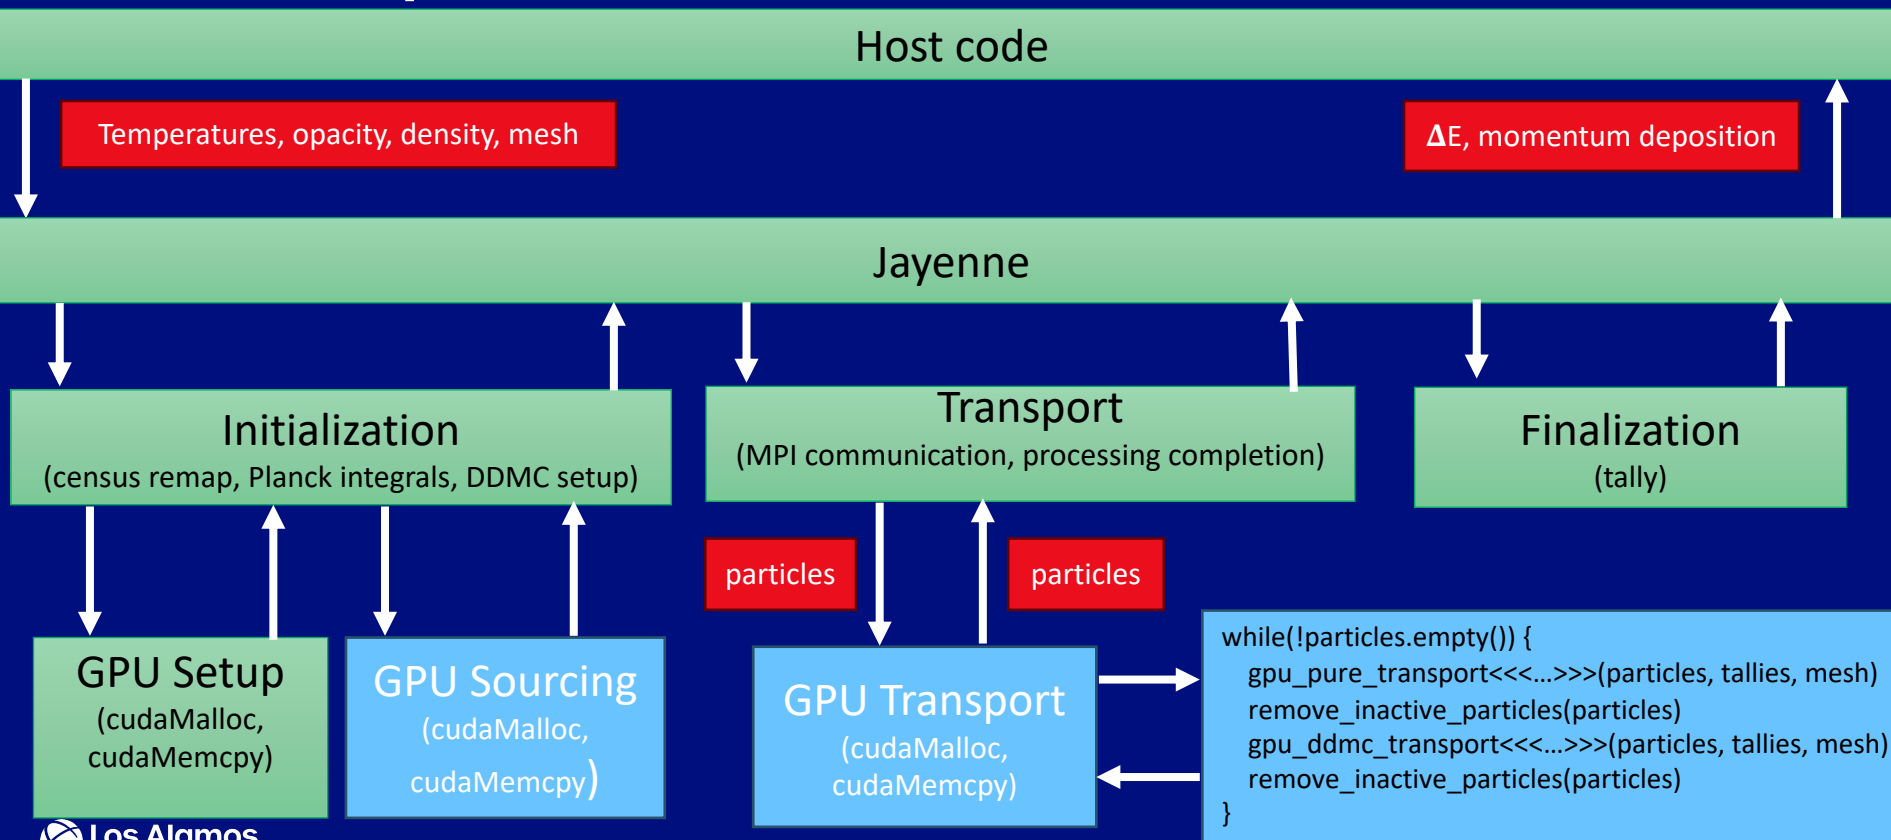
$$\frac{1}{c} \frac{\partial I}{\partial t} + \Omega \cdot \nabla I + \sigma_a I = \sigma_a B,$$
$$\frac{\partial U_m}{\partial t} = \int_0^\infty \int_0^{4\pi} \sigma_a I d\Omega d\nu - \int_0^\infty \sigma_a B d\nu + S_m.$$

- These equations are coupled by the absorption ( $\sigma_a I$ ) and emission terms ( $B$ )
- The emission term is proportional to  $T^4$ , making the equations nonlinear in  $T$
- The equations are linearized by assuming opacity, heat capacity and emission are fixed at the beginning of the timestep

# IMC linearizes the Thermal Radiative Transfer equations

- Implicit Monte Carlo introduces “effective scattering” to stabilize TRT
- A RHS term in energy and momentum equations of hydrodynamics
- Key points for computation:
  - Runs on the hydro mesh
  - Monte Carlo noise can seed hydrodynamic instabilities
  - Mean free paths can be very small (nanometers)
  - Diffusion acceleration methods are required for all practical problems
- IMC code at LANL is called Jayenne
  - 25 years old!
  - C++ with Fortran API to Cassio and Flag
  - Many improvements to standard Fleck and Cummings paper

# Rough outline of Jayenne code flow from host code to GPU transport



# Major refactoring coincided with GPU port

- Moving away from object-oriented design patterns at a low level
- Moving towards functional programming for tracking and sampling
- Focus on most-used features—2D AMR mesh, multigroup, basic tallies
- Dropping from four template arguments everywhere to one
- Simplify particle to 128 bytes and cell tallies to 88 bytes
- First targeted unaccelerated transport, then Random Walk acceleration, then Discrete Diffusion Monte Carlo (DDMC)

# Core design strategy

- GPU transport kernel templated on mesh type and dimension allows for a single compilation unit
- Explicit memory model and no allocations in device code
- One particle per thread, particles immediately loaded in to shared memory
- Truncated history, with particles sorted and requeued after 25 steps
- No virtual or recursive functions
- Hybrid history/event-based with two events: unaccelerated transport and Discrete Diffusion Monte Carlo (DDMC)
- CMake and simple preprocessor macros for mostly single-source (and HIP compilation)

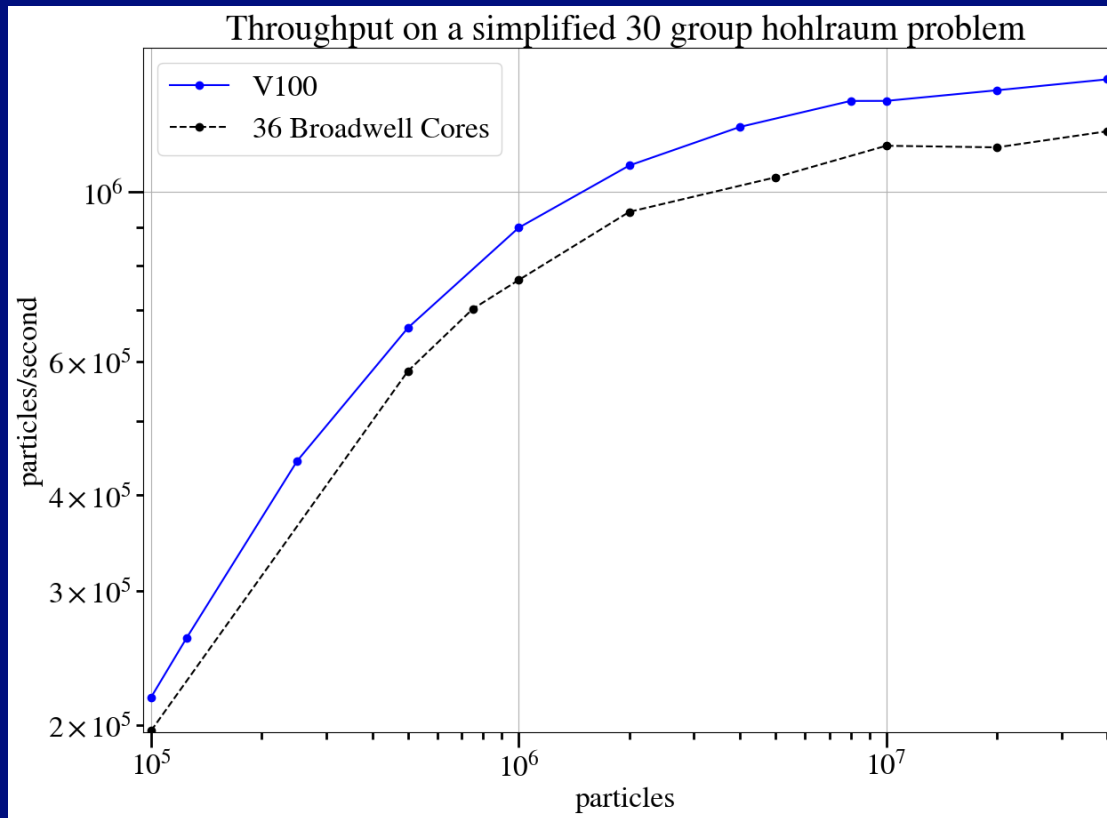


# Notes and oddities


- Most of our experience is on Sierra-like systems
- Better performance with GCC than XL
- CUDA 11.2+ showed significant slowdown (~2x)
- Best performance with CUDA's Multi-Process Service on
- Running with 40 ranks per node for best host code configuration
- Kernel sizes: No acceleration 113 registers, DDMC acceleration 96 registers
- Speedup from Broadwell nodes to both DDR and HBM Sapphire Rapids nodes is about 3.5x (roughly power scaling)

# Throughput saturates around 10 million particles

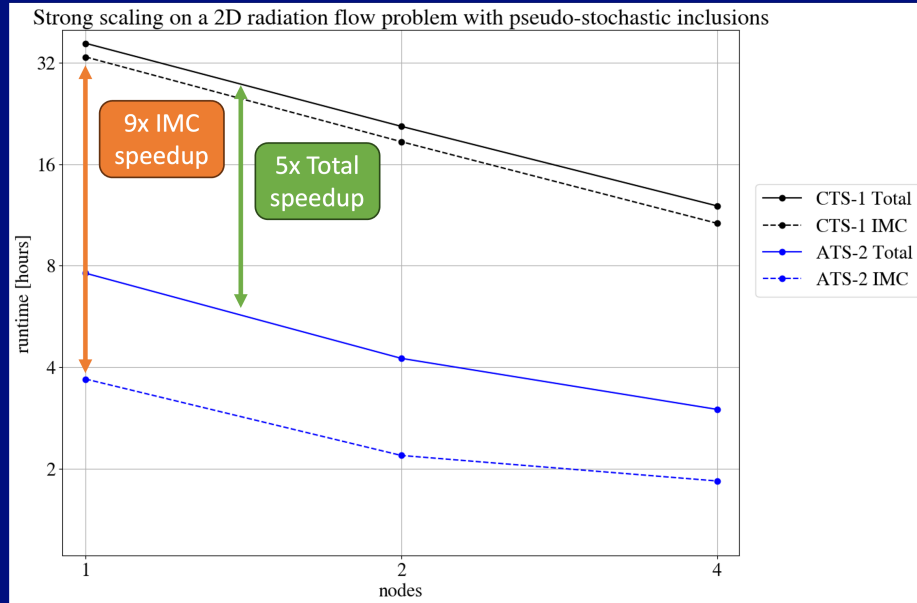
- Why doesn't CPU cross GPU?
- A bit odd looking, possibly due to relative expense of initialization with low particle counts
- Generally, expecting 4x speedup node-to-node



# Excellent speedup on smaller, transport dominated problems




Radiation temperature with an incident radiation source and background “inclusions”



At one node, 10 million particles per GPU (95% peak throughput)


# Good speedup on large, transport dominated problems

- 15 million particles per GPU at 8 nodes, about max throughput
- Runs out of memory on 4 nodes of CPU cluster
- Currently working with simple hydro mesh decomposition, moving towards internal decomposition and acquiescing MPI ranks



# Limited speedup on standard user runs

- Default cell counts gives 5 million particles per GPU at 8 nodes, 90% of peak throughput
- Does throughput curve take longer to saturate in domain-decomposed problems? How much longer?
- Way to measure this tail effect?



# Future work

- Port initialization routines
  - 45% time in transport at 80% of peak throughput
- Simplify momentum integral
  - Moving from dynamic quadrature points to many static points gives 20% speedup of GPU kernel
- GPU feature addition
  - External surface tallies
  - Compton scattering options
  - LD-IMC

# Conclusions and questions

- Trying to change how IMC is used at LANL by suggesting much more particles and fewer GPU nodes
- Is the GPU hiding MC memory latency? Does definitively knowing this matter for procurement or optimization work?
- I'm wary of event-based transport for IMC because particle can have many collisions within a timestep (10k kernel launches)
- After initialization, try a “persistent thread” approach
- Speedup above 8x is of limited utility without hydrodynamics and EOS on the GPU

# Compiling with -O2 and -res-usage on V100 with CUDA 11.1

```
ptxas info      : Compiling entry function '_ZN14rtt_imc_solver21hybrid_ddmc_transportIN12rtt_tracking20CylindricalCoord_sysILi2EEEN6rtt_mc16Cylindrical_MeshELi2EEEvPKNS1_9Cell_DataEPKNS13Cell_IMC_DataEPKNS14Cell_DDMC_DataIXT1_EEEPKNs1_2U3EPKdNS1_9FrequencyENS16Scattering_ModelEPNS10Cell_TallyEPNS8ParticleEjdPi' for 'sm_70'
ptxas info      : Function properties for '_ZN14rtt_imc_solver21hybrid_ddmc_transportIN12rtt_tracking20CylindricalCoord_sysILi2EEEN6rtt_mc16Cylindrical_MeshELi2EEEvPKNS1_9Cell_DataEPKNS13Cell_IMC_DataEPKNS14Cell_DDMC_DataIXT1_EEEPKNs1_2U3EPKdNS1_9FrequencyENS16Scattering_ModelEPNS10Cell_TallyEPNS8ParticleEjdPi'
2880 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Used 96 registers, 8192 bytes smem, 2496 bytes cmem[0], 432 bytes cmem[2]
ptxas info      : Function properties for '__internal_accurate_pow'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Function properties for '__internal_trig_reduction_slowpathd'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Compiling entry function '_ZN14rtt_imc_solver25hybrid_no_accel_transportIN12rtt_tracking20CylindricalCoord_sysILi2EEEN6rtt_mc16Cylindrical_MeshELi2EEEvPKNS1_9Cell_DataEPKNS13Cell_IMC_DataEPKNS14Cell_DDMC_DataIXT1_EEENS1_9FrequencyENS16Scattering_ModelEPKiPKNS1_5ArrayIdXmLT1_Li2EEEEPNs10Cell_TallyEPNS8ParticleEmdPid' for 'sm_70'
ptxas info      : Function properties for '_ZN14rtt_imc_solver25hybrid_no_accel_transportIN12rtt_tracking20CylindricalCoord_sysILi2EEEN6rtt_mc16Cylindrical_MeshELi2EEEvPKNS1_9Cell_DataEPKNS13Cell_IMC_DataEPKNS14Cell_DDMC_DataIXT1_EEENS1_9FrequencyENS16Scattering_ModelEPKiPKNS1_5ArrayIdXmLT1_Li2EEEEPNs10Cell_TallyEPNS8ParticleEmdPid'
2768 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Used 113 registers, 8192 bytes smem, 2504 bytes cmem[0], 480 bytes cmem[2]
ptxas info      : Function properties for '__internal_accurate_pow'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Function properties for '__internal_trig_reduction_slowpathd'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
```

- No acceleration transport uses 113 registers
- DDMC acceleration uses 96 registers